

RAISE-Arch: Refactoring Heterogeneous Systems Architectures for Scalable and Efficient AI Infrastructure

• Oct. 31, 2026 • Athens, Greece • Co-located with MICRO 2026

ABOUT THE WORKSHOP

RAISE-Arch is the first edition of an industry–academic workshop on architectural directions for scalable, efficient, and collaborative AI infrastructure. It addresses bottlenecks in computing resources, memory bandwidth and capacity, storage access, data movement, and interconnect scalability through technologies such as HBM, HBF, CXL, AI-SSD/computational storage, chiplets, heterogeneous accelerators, xPU computing, and tiered storage. Held in conjunction with [MICRO 2026](#), it brings together researchers and practitioners to discuss cross-layer solutions and hardware–software co-design and co-optimization for the Agentic AI era.

CALL FOR PAPERS

RAISE-Arch invites researchers, developers, and professionals from academia and industry to submit extended abstracts related to the workshop topics. Submissions may describe work-in-progress ideas, position and vision statements, benchmarking and evaluation studies, early-stage research, industry experience, or discussion-oriented contributions. The program will include invited talks and selected presentations based on the submitted contributions.

TOPICS OF INTEREST

- Heterogeneous architectures for new AI infrastructure
- Data- and memory-centric design: PIM, NDP, near-data, and in-memory computing
- Architectures for reducing data movement across compute, memory, and storage
- HBM, HBF, CXL, chiplets, and disaggregated-memory architectures for AI
- Computational storage, AI-SSD, and flash-assisted AI system architectures
- AI accelerators, NPUs, FPGAs, heterogeneous SoCs, and rack-scale AI platforms
- Redefining CPU architecture and capacity for Agentic AI infrastructure
- Benchmarking and evaluation methodologies for data-centric architectures
- Tracing and simulation methodologies for emerging workloads

SUBMISSION GUIDELINES

- Submissions in the form of extended abstracts related to the workshop topics.
- Submissions should be max 2 pages, excluding references, using a double-column format, and submitted as a printable PDF file.
- Each submission must be in English and include the title, author names, affiliations, and email. Submissions should not be anonymized.
- We encourage authors to use the provided LaTeX template: [Download LaTeX Template \(ZIP\)](#)
- No formal proceedings are planned, and submissions are intended to be non-archival.
- Presentation slides or related materials may be posted on the workshop website after the event, subject to author approval.
- At least one author of each accepted contribution is expected to attend the workshop and present the work.

IMPORTANT DATES

Paper submission	7th of Sept. 2026 (AoE)
Author notification	October 1, 2026 (AoE)
Workshop	October 31, 2026

WEBSITE & SUBMISSION



[WORKSHOP WEBSITE](#)



[SUBMIT A PAPER](#)

ORGANIZERS

- Andreas Koch (Technical University of Darmstadt)
- Sajjad Tamimi (Huawei Europe)
- Boris Grot (University of Edinburgh)
- Freddy Gabbay (Hebrew University of Jerusalem)
- Brian Hirano (XCENA)
- Zheng Li (Huawei Europe)
- Zhou Yigang (Huawei China)

PROGRAM COMMITTEE

- Mohammad Alian (Cornell University)
- Shaahin Angizi (New Jersey Institute of Technology)
- Ashkan Asgharzadeh (BSC)
- Stefano Cordo (Huawei)
- Jonas Dann (ETH Zürich)
- Sina Darabi (BSC)
- Holger Fröning (Heidelberg University)
- Jana Giceva (Technical University of Munich)
- João P. C. Lima (TU Dresden)
- Mohammad Loni (Volvo Autonomous Solutions)
- Haiyu Mao (King's College London)
- Geraldo F. Oliveira (Huawei)
- Daniele Passaretti (University Magdeburg)
- Ilia Petrov (Reutlingen University)
- Stratos Psomadakis (NTUA)
- Nezam Rohbani (BSC)
- Lukas Sommer (AMD)
- Leonardo Solis-Vasquez (TU Darmstadt)
- Nima TaheriNejad (Heidelberg University)
- Mohammad Kazem Taram (Purdue University)
- Sara Vinco (Politecnico di Torino)
- Chenyu Zeng (Huawei)

